PwC AGENTS WORKSHOP - HANDBOOK 2
Agent Patterns and Cost Reference

1. PROMPT AGENTS VERSUS HOSTED AGENTS

A prompt agent is declared entirely in configuration: a model, a set of instructions, and tools. The platform runs it. No container is built.

A hosted agent is your own code, packaged as a container image and run on Microsoft-managed infrastructure. Choose it when you need your own libraries, your own orchestration, or a custom protocol.

2. TOOLBOXES

A toolbox is a named bundle of tools shared across agents, with authentication and observability handled per tool call. Declaring tools in a toolbox rather than inline means several agents can reuse the same bundle and you can audit tool usage centrally.

3. MODEL ROUTER

model-router selects a model per request from a pool. Use it when request difficulty varies: simple turns go to a cheap model and hard turns escalate. Pin a specific model instead when you need output characteristics to be stable.

4. COST NOTES FOR THIS WORKSHOP

All labs use small serverless models on GlobalStandard SKU, billed per token with no reserved capacity. Hosted agents scale to zero when idle, so an agent nobody is talking to costs nothing to keep deployed.

The dominant cost in a class of this size is shared model throughput, not compute. Request a quota increase before the event.

5. MEMORY

A memory store extracts stable facts from conversations, consolidates them, and retrieves them on later turns. It needs a chat deployment and an embedding deployment. Without the Cognitive Services OpenAI User role, memory writes fail with 401 and the store stays silently empty.